

SparkulAkka: a Simple and Extensible Distributed Simulation System built on Spark and Akka

WSU Vancouver

Fabiana Ferracina

April 23, 2018

Overview

1 Introduction

- Motivation

2 Background

- Simulations
- The Actor Model
- RDDs

3 My Approach

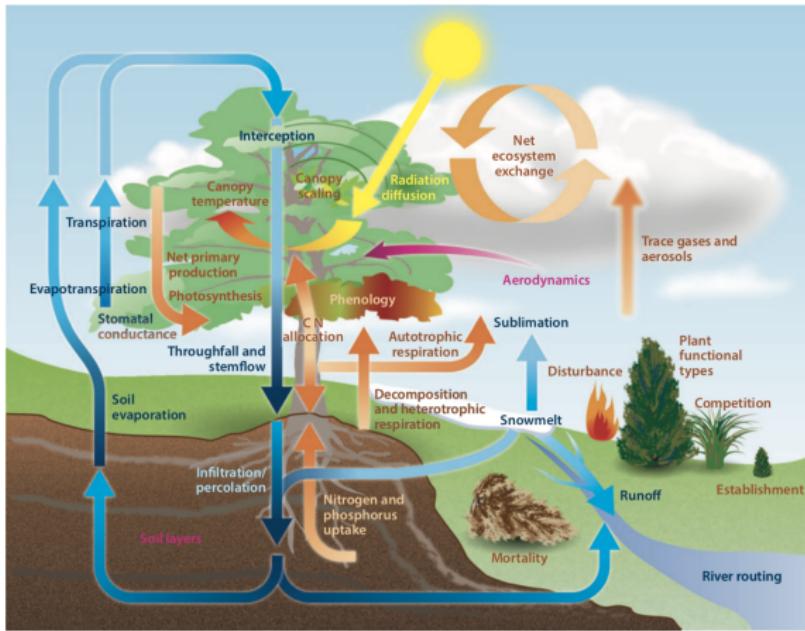
- Methods
- Results

4 The Future

- Goals and Aspirations

My research focus is on the modeling of forest dynamics and processes.

My research focus is on the modeling of forest dynamics and processes. My motivation is to use mathematics and computer science tools to model everything here:

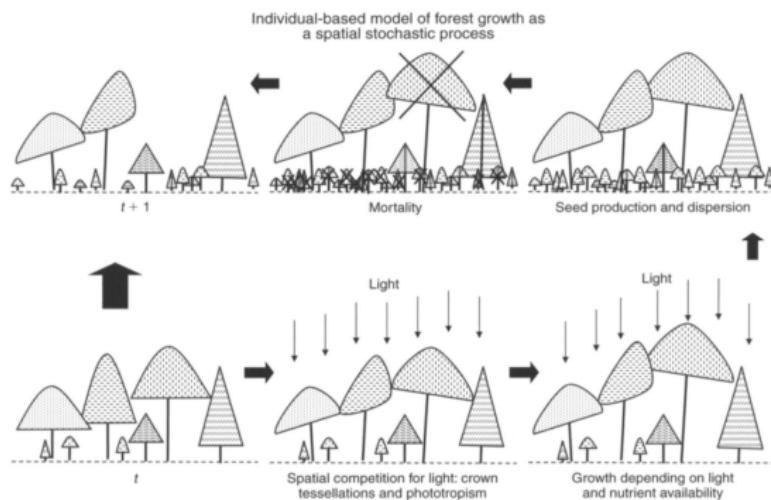


One step at a time...

Let's focus on succession dynamics for now:

One step at a time...

Let's focus on succession dynamics for now:



Many Questions!

- How do we implement sustainable forestry?

Many Questions!

- How do we implement sustainable forestry?
- How long will it take for a forest to recover after a large fire?

Many Questions!

- How do we implement sustainable forestry?
- How long will it take for a forest to recover after a large fire?
- What will happen to the Northwest forests if temperatures increase by 2 degrees Celsius?

Many Questions!

- How do we implement sustainable forestry?
- How long will it take for a forest to recover after a large fire?
- What will happen to the Northwest forests if temperatures increase by 2 degrees Celsius?

Data can help, also a computer or 50...



Overview

1 Introduction

- Motivation

2 Background

- Simulations
- The Actor Model
- RDDs

3 My Approach

- Methods
- Results

4 The Future

- Goals and Aspirations

Simulations are computed scenarios based on a simplified version of reality usually captured by differential equations.

Simulations are computed scenarios based on a simplified version of reality usually captured by differential equations.

Natural processes are usually stochastic and non-linear.

Simulations are computed scenarios based on a simplified version of reality usually captured by differential equations.

Natural processes are usually stochastic and non-linear.

We want to compute new values, evolve the system given these values, and loop for thousands of cycles over a large area containing millions of entities

Simulations are computed scenarios based on a simplified version of reality usually captured by differential equations.

Natural processes are usually stochastic and non-linear.

We want to compute new values, evolve the system given these values, and loop for thousands of cycles over a large area containing millions of entities

... and we want to use previous data and keep history...

Simulations are computed scenarios based on a simplified version of reality usually captured by differential equations.

Natural processes are usually stochastic and non-linear.

We want to compute new values, evolve the system given these values, and loop for thousands of cycles over a large area containing millions of entities

... and we want to use previous data and keep history...

Interestingly, active forest modeling research, despite the many mentions of computational complexity and data-intensiveness has a gap when it comes to distributed computing.

Overview

1 Introduction

- Motivation

2 Background

- Simulations
- The Actor Model
- RDDs

3 My Approach

- Methods
- Results

4 The Future

- Goals and Aspirations

- An approach to understand and reason about concurrency.

- An approach to understand and reason about concurrency.
- Executing objects, called Actors, communicate exclusively through asynchronous messages.

- An approach to understand and reason about concurrency.
- Executing objects, called Actors, communicate exclusively through asynchronous messages.
- State is maintained privately within each Actor, thus deadlocks and data integrity pose little to no issues in an Actor-oriented system.

Overview

1 Introduction

- Motivation

2 Background

- Simulations
- The Actor Model
- RDDs

3 My Approach

- Methods
- Results

4 The Future

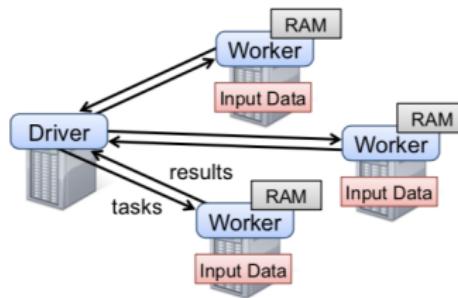
- Goals and Aspirations

- Spark - an in-memory distributed data analysis framework.

- Spark - an in-memory distributed data analysis framework.
- RDDs: Resilient Distributed Datasets - a distributed memory abstraction.

- Spark - an in-memory distributed data analysis framework.
- RDDs: Resilient Distributed Datasets - a distributed memory abstraction.
- RDDs allow for fault-tolerant computations on large clusters by providing a restricted form of shared memory based on coarse-grained transformations.

- Spark - an in-memory distributed data analysis framework.
- RDDs: Resilient Distributed Datasets - a distributed memory abstraction.
- RDDs allow for fault-tolerant computations on large clusters by providing a restricted form of shared memory based on coarse-grained transformations.
- Spark Runtime:



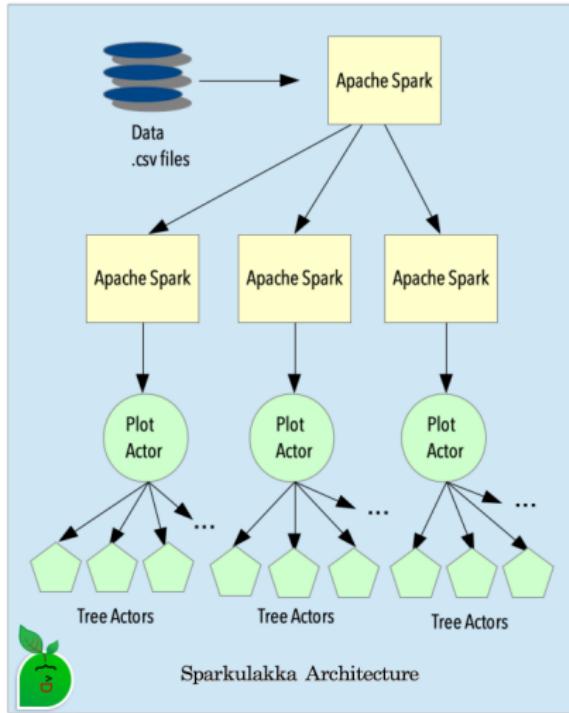


Figure: SparkulAkka architecture and a Harvard Forest flux tower

Overview

1 Introduction

- Motivation

2 Background

- Simulations
- The Actor Model
- RDDs

3 My Approach

- Methods
- Results

4 The Future

- Goals and Aspirations

- Spark to load, process and analyze forest measurement data -
Sparkulation piece where further statistical models can be
adapted.

- Spark to load, process and analyze forest measurement data - Sparkulation piece where further statistical models can be adapted.
- Data sources: FIADB, WorldClim, Harvard Forest flux towers

- Spark to load, process and analyze forest measurement data - Sparkulation piece where further statistical models can be adapted.
- Data sources: FIADB, WorldClim, Harvard Forest flux towers
- Sample Spark SQL:

```
SELECT plot, spcd, AVG(ht), AVG(dia)
FROM IL_TREE_csv
WHERE statuscd = 1 AND diahtcd = 1 AND ht IS NOT NULL
      AND invyr = "2008"
GROUP BY 1, 2 ORDER BY 1, 2;
```

- Spark to load, process and analyze forest measurement data - Sparkulation piece where further statistical models can be adapted.
- Data sources: FIADB, WorldClim, Harvard Forest flux towers
- Sample Spark SQL:

```
SELECT plot, spcd, AVG(ht), AVG(dia)
FROM IL_TREE_csv
WHERE statuscd = 1 AND diahtcd = 1 AND ht IS NOT NULL
      AND invyr = "2008"
GROUP BY 1, 2 ORDER BY 1, 2;
```

- WHERE conditions are filter operations. GROUP BY is a groupByKey (with reduce for calculating the final values via AVG). ORDER BY is a sort.

- Initial differential model variables are then fed into a simulation using the Actor Model framework Akka.

- Initial differential model variables are then fed into a simulation using the Actor Model framework Akka.
- As we can see Actors fully utilize available cores:

Processes: 295 total, 3 running, 292 sleeping, 1216 threads 20:55:53
Load Avg: 4.79, 2.26, 1.89 CPU usage: 98.28% user, 1.71% sys, 0.0% idle SharedLibs: 213M resident, 63M data, 26M linkedit.
MemRegions: 26043 total, 2365M resident, 165M private, 523M shared. PhysMem: 7747M used (1466M wired), 444M unused.
VM: 1321G vszsize, 1297M framework vszsize, 0(0) swapins, 0(0) swapouts. Networks: packets: 84178/42M in, 67523/12M out.
Disks: 245300/6196M read, 133444/4176M written.

PID	COMMAND	%CPU	TIME	#TH	#WQ	#PORT	MEM	PURG	CMPRS	PGRP	PPID	STATE	BOOSTS	%CPU_ME	%CPU_OTHRS	UID	FAULTS
922	screenCaptur	0.1	00:00:16	4	2	152	3712K	0B	0B	922	1	sleeping	*0[1634+]	0.00000	0.04138	501	5498
921	screenCaptur	1.2	00:01:13	2	1	54	3304K	636K	0B	308	308	sleeping	*0[1]	0.04138	0.00000	501	112271+
920	java	392.1	01:03:51	41/8	1	127	682M+	0B	0B	453	453	running	*0[1]	0.00000	0.00000	501	187405+
919	java	8.0	00:03:06	29	1	103	144M	0B	0B	453	453	sleeping	*0[1]	0.00000	0.00000	501	42976
918	java	8.0	00:01:45	33	1	111	74M	0B	0B	453	1	sleeping	*0[1]	0.00000	0.00000	501	24572
917	ocspd	8.0	00:00:02	2	1	27	1340K	0B	0B	917	1	sleeping	*0[1]	0.00000	0.00000	0	2131
916	java	0.0	00:09:66	31	1	107	198M	0B	0B	453	453	sleeping	*0[1]	0.00000	0.00000	501	91106
911	mdworker_sha	8.0	00:00:04	3	1	59	3072K	0B	0B	911	1	sleeping	*0[1]	0.00000	0.00000	501	3778
910	mdworker_sha	8.0	00:00:04	3	1	59	3112K	0B	0B	910	1	sleeping	*0[1]	0.00000	0.00000	501	3784
901	java	0.0	00:00:02	2	1	15	3706K	0B	0B	901	1	sleeping	*0[1]	0.00000	0.00000	501	1248

- Initial differential model variables are then fed into a simulation using the Actor Model framework Akka.
- As we can see Actors fully utilize available cores:

```
Processes: 295 total, 3 running, 292 sleeping, 1216 threads          20:55:53
Load Avg: 4.79, 2.26, 1.89 CPU usage: 98.28% user, 1.71% sys, 0.0% idle SharedLibs: 213M resident, 63M data, 26M linkedit.
MemRegions: 26043 total, 2365M resident, 165M private, 523M shared. PhysMem: 7747M used (1466M wired), 444M unused.
VM: 1321G vszie, 1297M framework vszie, 0(0) swapins, 0(0) swapouts. Networks: packets: 84178/42M in, 67523/12M out.
Disks: 245300/6196M read, 133444/4176M written.

          12.107692307692307

```

PID	COMMAND	%CPU	TIME	#TH	#WQ	#PORT	MEM	PURG	CMPRS	PGRP	PPID	STATE	BOOSTS	%CPU_ME	%CPU_OTHRS	UID	FAULTS
922	screenCaptur	0.1	00:00:16	4	2	152	3712K	0B	0B	922	1	sleeping	*@{1634+}	0.00000	0.04138	501	5498
921	screenCaptur	1.2	00:01:13	2	1	54	3304K	636K	0B	308	308	sleeping	*@{1}	0.04138	0.00000	501	112271+
920	java	392.1	01:03:51	41/8	1	127	682M+	0B	0B	453	453	running	*@{1}	0.00000	0.00000	501	187405+
919	java	8.0	00:03:06	29	1	103	144M	0B	0B	453	453	sleeping	*@{1}	0.00000	0.00000	501	42976
918	java	8.0	00:01:45	33	1	111	74M	0B	0B	453	1	sleeping	*@{1}	0.00000	0.00000	501	24572
917	ocspd	8.0	00:00:02	2	1	27	1340K	0B	0B	917	1	sleeping	*@{1}	0.00000	0.00000	0	2131
916	java	0.0	00:09:06	31	1	107	198M	0B	0B	453	453	sleeping	*@{1}	0.00000	0.00000	501	91106
911	mdworker_sha	8.0	00:00:04	3	1	59	3072K	0B	0B	911	1	sleeping	*@{1}	0.00000	0.00000	501	3778
910	mdworker_sha	8.0	00:00:04	3	1	59	3112K	0B	0B	910	1	sleeping	*@{1}	0.00000	0.00000	501	3784
904	java	0.0	00:00:02	2	1	15	3766K	0B	0B	904	1	sleeping	*@{1}	0.00000	0.00000	501	1242

- Whereas the sequential implementation does not:

```
Processes: 279 total, 3 running, 276 sleeping, 1098 threads          19:25:59
Load Avg: 1.56, 1.58, 1.71 CPU usage: 25.80% user, 2.7% sys, 72.11% idle SharedLibs: 288M resident, 63M data, 37M linkedit.
MemRegions: 25359 total, 2449M resident, 158M private, 446M shared. PhysMem: 7362M used (1424M wired), 829M unused.
VM: 1254G vszie, 1297M framework vszie, 0(0) swapins, 0(0) swapouts. Networks: packets: 22615/6950K in, 19328/4691K out.
Disks: 193532/4866M read, 91764/3399M written.

          14.1708997 18997189873

```

PID	COMMAND	%CPU	TIME	#TH	#WQ	#PORT	MEM	PURG	CMPRS	PGRP	PPID	STATE	BOOSTS	%CPU_ME	%CPU_OTHRS	UID	FAULTS	
666	screenCaptur	0.0	00:00:31	3	1	148	3576K	0B	0B	666	1	sleeping	*@{3113}	0.00000	0.00000	501	5515	
665	screenCaptur	0.0	00:02:71	2	1	56	3360K	636K	0B	308	308	sleeping	*@{1}	0.00000	0.00000	501	276900	
664	Python	99.5	00:41:89	1/1	0	10	15M	0B	0B	664	632	running	*@{1}	0.00000	0.00000	501	5093	
650	com.apple.1C	0.0	00:00:05	2	1	55	3072K	0B	0B	650	1	sleeping	*@{1}	0.00000	0.00000	501	4396	
632	bash	simul0.0	00:00:04	1	1	10000	21	892K	0B	0B	632	631	sleeping	*@{1}	0.00000	0.00000	501	1524
631	login	0.0	00:00:01	2	1	31	1032K	0B	0B	631	488	sleeping	*@{9}	0.00000	0.00000	0	1741	
630	mdworker_sha	0.0	00:00:07	4	1	61	4580K	0B	0B	630	1	sleeping	*@{1}	0.00000	0.00000	501	8055	
624	trustd	0.0	00:00:09	2	10	37	4188K	168K	0B	624	1	sleeping	*@{35}	0.00000	0.00000	89	3012	

- Initial differential model variables are then fed into a simulation using the Actor Model framework Akka.
- As we can see Actors fully utilize available cores:

Processes: 295 total, 3 running, 292 sleeping, 1216 threads														20:55:53				
PID	COMMAND	%CPU	TIME	#TH	#WQ	#PORT	MEM	PURG	CMPRS	PGRP	PPID	STATE	BOOSTS	%CPU_ME	%CPU_OTHRHS	UID	FAULTS	
922	screenCaptur	0.1	00:00:16	4	2	152	3712K	0B	0B	922	1	sleeping	*[8 1634+]	0.00000	0.04138	501	5498	
921	screenCaptur	1.2	00:01:13	2	1	54	3304K	636K	0B	308	308	sleeping	*[0 1]	0.04138	0.00000	501	112271+	
920	java	392.1	01:03:51	41/8	1	127	682M+	0B	0B	453	453	running	*[0 1]	0.00000	0.00000	501	187405+	
919	java	8.0	00:03:06	29	1	103	144M	0B	0B	453	453	sleeping	*[0 1]	0.00000	0.00000	501	42976	
918	java	8.0	00:01:45	33	1	111	74M	0B	0B	453	1	sleeping	*[0 1]	0.00000	0.00000	501	24572	
917	ocspd	8.0	00:00:00	2	2	27	1340K	0B	0B	917	1	sleeping	*[0 1]	0.00000	0.00000	0	2131	
916	java	0.0	00:09:06	31	1	107	198M	0B	0B	453	453	sleeping	*[0 1]	0.00000	0.00000	501	91106	
911	mdworker_sha	0.0	00:00:00	4	3	1	59	3072K	0B	0B	911	1	sleeping	*[0 1]	0.00000	0.00000	501	3778
910	mdworker_sha	0.0	00:00:00	4	3	1	59	3112K	0B	0B	910	1	sleeping	*[0 1]	0.00000	0.00000	501	3784
904	java	0.0	00:00:00	2	2	15	3760K	0B	0B	904	1	sleeping	*[0 1]	0.00000	0.00000	501	1242	

- Whereas the sequential implementation does not:

Processes: 279 total, 3 running, 276 sleeping, 1098 threads														19:25:59				
PID	COMMAND	%CPU	TIME	#TH	#WQ	#PORT	MEM	PURG	CMPRS	PGRP	PPID	STATE	BOOSTS	%CPU_ME	%CPU_OTHRHS	UID	FAULTS	
666	screenCaptur	0.0	00:00:31	3	1	148	3576K	0B	0B	666	1	sleeping	*[0 313]	0.00000	0.00000	501	5515	
665	screenCaptur	0.0	00:02:71	2	1	56	3360K	636K	0B	308	308	sleeping	*[0 1]	0.00000	0.00000	501	276900	
664	Python	99.5	00:41:89	1/1	0	10	15M	0B	0B	664	632	running	*[0 1]	0.00000	0.00000	501	5093	
650	com.apple.1C	0.0	00:00:05	2	1	55	3072K	0B	0B	650	1	sleeping	*[0 1]	0.00000	0.00000	501	4396	
632	bash	simul0.0	00:00:04	1	1	1000	21	892K	0B	0B	632	631	sleeping	*[0 1]	0.00000	0.00000	501	1524
631	login	0.0	00:00:01	2	1	31	1032K	0B	0B	631	488	sleeping	*[0 9]	0.00000	0.00000	0	1741	
630	mdworker_sha	0.0	00:00:07	4	2	61	4580K	0B	0B	630	1	sleeping	*[0 1]	0.00000	0.00000	501	8055	
624	trustd	0.0	00:00:09	2	2	37	4188K	168K	0B	624	1	sleeping	*[0 35]	0.00000	0.00000	89	3012	

- SimulAkka thus has promise of horizontal scalability.

Systems used

- Sparkulation was run via Databricks.com hosting, which provides a cluster with a total of 6GB of memory - this ran all the code snippets that interacted with the underlying Spark framework.

Systems used

- Sparkulation was run via Databricks.com hosting, which provides a cluster with a total of 6GB of memory - this ran all the code snippets that interacted with the underlying Spark framework.
- SimulAkka was run on a system with a quad-core Intel i5 3.4 GHz processor and 8GB of memory - this ran both the SimulAkka-powered simulation, and a simpler (non-Actor model) Python implementation for comparison.

Overview

1 Introduction

- Motivation

2 Background

- Simulations
- The Actor Model
- RDDs

3 My Approach

- Methods
- Results

4 The Future

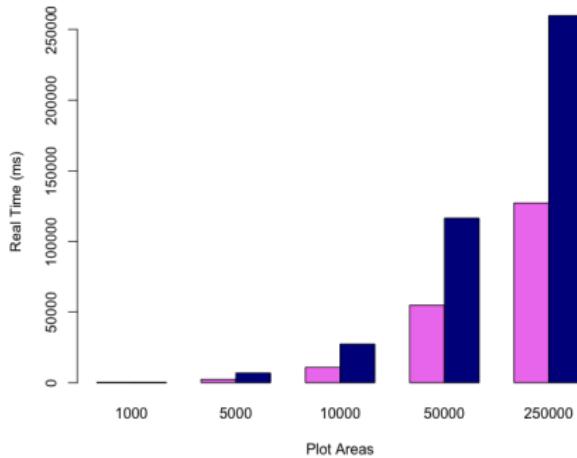
- Goals and Aspirations

Sparkulation performance (versus R)

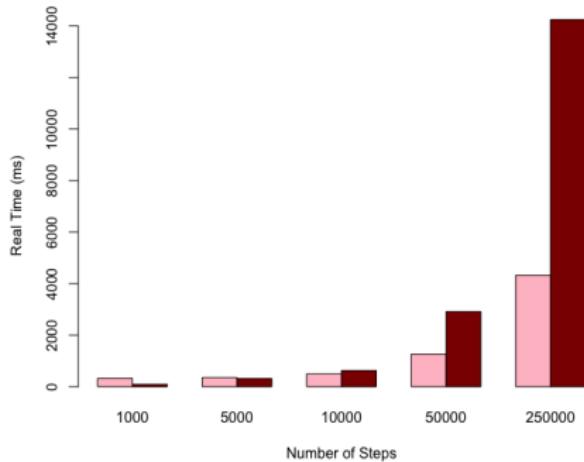
- Spark was 20% faster than R in my trial run of a single, relatively small (73mb) CSV that easily fits in most home computer's memory. Processing the entirety (21 gigabytes+) of FIADB data would give a significant advantage to Spark, both in feasibility and performance.

SimulAkka performance

Akka (violet) v. Python (blue)



Akka (pink) v. Python (red)



Overview

1 Introduction

- Motivation

2 Background

- Simulations
- The Actor Model
- RDDs

3 My Approach

- Methods
- Results

4 The Future

- Goals and Aspirations

```
void Plant() {  
    //#pragma omp parallel for num_threads(NUM_THREADS)  
    for (unsigned short s = 0; s < nS; s++) {  
        cohort[K].s = s;  
        //            if(OH==0) value[s].caTotal=FA/2;  
        //            cohort[K].n = (species[s].n /(double)nP) *  
        // (value[s].caTotal/FA);  
        cohort[K].n = species[s].n / (double)nP;  
        cohort[K].d = species[s].d;  
        cohort[K].h = height(cohort[K]);  
        cohort[K].r = radius(cohort[K]);  
        K++;  
        value[s].kTotal++;  
        value[nS].kTotal++;  
    }  
}  
...
```

- Automatically link the Spark pipeline with the Akka simulation system

- Automatically link the Spark pipeline with the Akka simulation system
- Fully Implement Sortie-PPA in SimulAkka

- Automatically link the Spark pipeline with the Akka simulation system
- Fully Implement Sortie-PPA in SimulAkka
- Explore the deployment of Spark on Amazon EMR to process over 20 million trees

- Automatically link the Spark pipeline with the Akka simulation system
- Fully Implement Sortie-PPA in SimulAkka
- Explore the deployment of Spark on Amazon EMR to process over 20 million trees
- Add soil biogeochemistry (C and N) to SimulAkka

- Automatically link the Spark pipeline with the Akka simulation system
- Fully Implement Sortie-PPA in SimulAkka
- Explore the deployment of Spark on Amazon EMR to process over 20 million trees
- Add soil biogeochemistry (C and N) to SimulAkka
- Port climate change models to SimulAkka and add bioclimatic data processing to Sparkulation

- Automatically link the Spark pipeline with the Akka simulation system
- Fully Implement Sortie-PPA in SimulAkka
- Explore the deployment of Spark on Amazon EMR to process over 20 million trees
- Add soil biogeochemistry (C and N) to SimulAkka
- Port climate change models to SimulAkka and add bioclimatic data processing to Sparkulation
- Once fast data analysis pipeline is ready, reparametrize and update models

- Automatically link the Spark pipeline with the Akka simulation system
- Fully Implement Sortie-PPA in SimulAkka
- Explore the deployment of Spark on Amazon EMR to process over 20 million trees
- Add soil biogeochemistry (C and N) to SimulAkka
- Port climate change models to SimulAkka and add bioclimatic data processing to Sparkulation
- Once fast data analysis pipeline is ready, reparametrize and update models
- Make attractive and healthy ecology codebase in order to encourage code reuse and maintenance by scientific community

Questions?

